

Machine Learning- Exercise 3

SVM, AdaBoost

George Lydakis & Idil Esen Zulfikar
`<lastname>@vision.rwth-aachen.de`

RWTH Aachen University - Computer Vision Group
<http://www.vision.rwth-aachen.de/>

2022-12-08

Content

1. Support Vector Machine

2. AdaBoost

Content

1. Support Vector Machine

2. AdaBoost

Support Vector Machine - Recap

- ▶ The SVM tries to find a classifier which maximizes the margin between 2 classes.
- ▶ Up to now considered linear classifiers

$$y(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b$$

- ▶ Formulation as a complex optimization problem
- ▶ Find the hyperplane satisfying

$$\operatorname{argmin}_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{W}\|^2$$

- ▶ under the constraints

$$t_n(\mathbf{w}^T \mathbf{x}_n + b) \geq 1, \forall n$$

- ▶ based on training data points \mathbf{x}_n and target value $t_n \in \{-1, 1\}$

Support Vector Machine - Recap

- Lagrangian primal form

$$\begin{aligned}\mathcal{L}_p &= \frac{1}{2} \|\mathbf{w}\|^2 - \sum_{n=1}^N \alpha_n \{t_n(\mathbf{w}^T \mathbf{x}_n + b) - 1\} \\ &= \frac{1}{2} \|\mathbf{w}\|^2 - \sum_{n=1}^N \alpha_n \{t_n y(\mathbf{x}_n) - 1\}\end{aligned}$$

- The solution of \mathcal{L}_p needs to fulfill the KKT conditions
 - Necessary and sufficient conditions

$$\begin{array}{ll}\alpha_n \geq 0 & \lambda \geq 0 \\ t_n y(\mathbf{x}_n) - 1 \geq 0 & f(\mathbf{x}) \geq 0 \\ \alpha_n \{t_n y(\mathbf{x}_n) - 1\} = 0 & \lambda f(\mathbf{x}) = 0\end{array}$$

Support Vector Machine - Recap

- ▶ Solution for hyperplane
 - ▶ computed as linear combination of the training examples


$$\mathbf{w} = \sum_{n=1}^N \alpha_n t_n \mathbf{x}_n$$

- ▶ Suppose solution: $\alpha_n \neq 0$ only for some points, **the support vectors**
 - ▶ Only the SVs actually influence the decision Boundary!
- ▶ Compute b by averaging over all support vectors:

$$b = \frac{1}{N_S} \sum_{n \in S} \left(t_n - \sum_{m \in S} \alpha_m t_m \mathbf{x}_m^T \mathbf{x}_n \right)$$

Support Vector Machine - Recap

- ▶ The training points for which $\alpha_n \geq 0$ are called support vectors



- ▶ Graphical interpretation:
 - ▶ The support vectors are the points on the margin.
 - ▶ They define the margin and thus the hyperplane.
 - ▶ All other points can be discarded.

Support Vector Machine - Recap

- ▶ Primal form has time complexity of $\mathcal{O}(D^3)$ in D dimensions
- ▶ Dual form can be obtained by substituting weight vectors

$$\mathbf{w} = \sum_{n=1}^N \alpha_n t_n \mathbf{x}_n$$

$$\mathcal{L}_d(\alpha) = \sum_{n=1}^N \alpha_n - \frac{1}{2} \sum_{n=1}^N \sum_{m=1}^N \alpha_n \alpha_m t_n t_m (\mathbf{x}_m^T \mathbf{x}_n)$$

- ▶ Under the conditions


$$\alpha_n \geq 0 \quad \forall n$$

$$\sum_{n=1}^N \alpha_n t_n = 0$$

- ▶ time complexity between $\mathcal{O}(N)$ and $\mathcal{O}(N^2)$

Support Vector Machine - Recap

- In case of overlapping class distribution, we allow some points to be miss-classified but with linear penalty that increases with the distance from the boundary.



Support Vector Machine - Recap

- ▶ Slack variables
 - ▶ $\xi \geq 0$ for each data point.
- ▶ Interpretation
 - ▶ $\xi_n = |t_n - y(\mathbf{x}_n)|$
 - ▶ $\xi_n = 0$ for points correctly classified or on the margin
 - ▶ $0 \leq \xi_n \leq 1$ points lie inside the margin on the correct side of the decision boundary
 - ▶ $y(\mathbf{x}_n) = 0$ and $\xi = 1$ for the data point on the decision boundary
 - ▶ $\xi_n > 1$ for misclassified points
- ▶ The exact classification constraints are then replaced with

$$t_n y(\mathbf{x}) \geq 1 - \xi_n$$

Support Vector Machine - Recap

- ▶ Goal is to maximize margin while softly penalizing points that lie on the wrong side of the margin boundary.
- ▶ Therefore we minimize :

$$\mathcal{C} \sum_{n=1}^N \xi_n + \frac{1}{2} \|\mathbf{w}\|^2$$

- ▶ Where the parameter \mathcal{C} controls the trade-of between the slack variable penalty and the margin.
- ▶ In the limit of $\mathcal{C} \rightarrow \infty$ we will recover the earlier, fully linearly separable case (non overlapping class distribution).

Support Vector Machine - Recap

- ▶ The corresponding Lagrangian multiplier can be given:

$$\mathcal{L}(\mathbf{w}, b, \alpha, \xi) = \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{n=1}^N \xi_n - \sum_{n=1}^N \alpha_n \{t_n y(\mathbf{x}_n) - 1 + \xi_n\} - \sum_{n=1}^N \mu_n \xi_n$$

- ▶ Where α_n and μ_n are Lagrangian multipliers
- ▶ The corresponding set of KKT conditions are given by:

$$\begin{aligned} \alpha_n &\geq 0 & \mu_n &\geq 0 \\ t_n y(\mathbf{x}_n) - 1 + \xi_n &\geq 0 & \xi_n &\geq 0 \\ \alpha_n \{t_n y(\mathbf{x}_n) - 1 + \xi_n\} &= 0 & \mu_n \xi_n &= 0 \end{aligned}$$

- ▶ We can optimize $\mathbf{w}, b, \{\xi_n \mid n \in \{1, \dots, N\}\}$ as follows:

$$\begin{aligned} \frac{\partial \mathcal{L}}{\partial \mathbf{w}} = 0 & \Rightarrow \mathbf{w} = \sum_{n=1}^N \alpha_n t_n \mathbf{x}_n \\ \frac{\partial \mathcal{L}}{\partial b} = 0 & \Rightarrow 0 = \sum_{n=1}^N \alpha_n t_n \\ \frac{\partial \mathcal{L}}{\partial \xi_n} = 0 & \Rightarrow a_n = C - \mu_n \end{aligned}$$

Support Vector Machine - Recap

- New SVM Dual: Maximize

$$\mathcal{L}_d(\alpha) = \sum_{n=1}^N \alpha_n - \frac{1}{2} \sum_{n=1}^N \sum_{m=1}^N \alpha_n \alpha_m t_n t_m (\mathbf{x}_m^T \mathbf{x}_n)$$

- Under conditions

$$0 \leq \alpha_n \leq \mathcal{C}$$

$$\sum_{n=1}^N \alpha_n t_n = 0$$

- This is a quadratic programming problem

Support Vector Machine - Recap

- ▶ Using one of the KKT condition and result from partial derivation of Lagrange function we can derive condition for data points to be either support vector or slacks as below,
- ▶ We know that

$$\frac{\partial L}{\partial \xi_n} = 0 \quad \Rightarrow \quad \alpha_n = C - \mu_n$$

and

$$\mu_n \geq 0$$

$$\xi_n \geq 0$$

$$\mu_n \xi_n = 0$$

- ▶ For support vectors and slacks $0 \leq \alpha_n \leq C$
- ▶ Moreover for slacks $\xi_n \geq 0$ which implies $\mu_n = 0$
- ▶ From the partial derivation shown above $\rightarrow \alpha_n = C$

Content

1. Support Vector Machine


2. AdaBoost

Adaboost[Freund & Schapire, 1996]-Recap

- ▶ Main idea
 - ▶ Instead of resampling, reweight misclassified training examples.
 - ▶ Increase the chance of being selected in a sampled training set.
 - ▶ Or increase the misclassification cost when training on the full set.
- ▶ Components
 - ▶ $c_k(\mathbf{x})$: “weak” or base classifier
 - ▶ Condition: $< 50\%$ training error over any distribution
 - ▶ $C(\mathbf{x})$: “strong or final classifier”
- ▶ Adaboost:
 - ▶ Construct a strong classifier as a thresholded linear combination of the weighted classifiers:

$$C(\mathbf{x}) = \text{sign} \left(\sum_{k=1}^K \alpha_k c_k(\mathbf{x}) \right)$$

Adaboost - Recap




Consider a 2D feature space with **positive** and **negative** examples.

Each weak classifier splits the training examples with at least 50% accuracy.

Examples misclassified by a previous weak learner are given more emphasis at future rounds.

Adaboost - Recap



- Final classifier is combination of the weak classifiers.

Adaboost - Algorithm

- ▶ Initialization: Set $w_n^{(1)} = \frac{1}{N}$ for $n = 1, \dots, N$.
- ▶ For $k = 1, \dots, k$ iterations
 - ▶ Train a new weak classifier $c_k(\mathbf{X})$ using current weights $\mathbf{W}^{(k)}$ by minimizing the weighted error function
 - ▶ estimate the weighted error of this classifier on \mathbf{X} :

$$\epsilon_k = \frac{\sum_{n=1}^N w_n^{(k)} I(c_k(\mathbf{X}) \neq y_n)}{\sum_{n=1}^N w_n^{(k)}}$$

- ▶ Calculate a weighting coefficient for $c_k(\mathbf{X})$:

$$\alpha_k = \ln \left\{ \frac{1 - \epsilon_k}{\epsilon_k} \right\}$$

- ▶ Update the weighting coefficients:

$$w_n^{(k+1)} = w_n^{(k)} \exp\{\alpha_k I(c_k(\mathbf{X}) \neq y_n)\}$$